

# Spatial filtering

Semester 2, 2022

Kris Ehinger







# Outline

- Introduction to convolution
- Commonly-used linear filters
- Filters in practice

# Learning outcomes

- Explain the convolution and cross-correlation operations
- Identify commonly-used filters and their expected outputs
- Explain practical considerations in implementing filters (efficiency, border handling)

# Introduction to convolution

# Pixel operator


- Pixel operator: computes an output value at each pixel location, based on the input pixel value

$$g(i, j) = h(f(i, j))$$

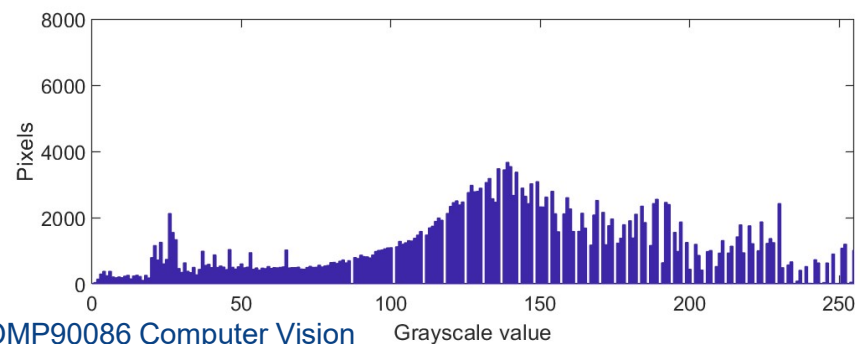


Output image  $g$

Input image  $f$

- Example:  $g(i, j) = 0.5(f(i, j))$



# Gamma correction



# Local operator

- Local operator: computes an output value at each pixel location, based on a neighbourhood of pixels around the input pixel
- Example: sharpening filter





# Linear filtering

- Output pixel's value is a weighted sum of a neighbourhood around the input pixel

$$g(i, j) = h(u, v) \otimes f(i, j)$$

Output image  $g$       Kernel  $h$       Input image  $f$


Cross-correlation convolution

$$g(i, j) = \sum_{u, v} f(i + u, j + v) h(u, v)$$

# Linear filtering

- Output pixel's value is a weighted sum of a neighbourhood around the input pixel

$$g(i, j) = h(u, v) * f(i, j)$$

Output image  $g$       Kernel  $h$        Input image  $f$

Convolution operator

$$g(i, j) = \sum_{u, v} f(i - u, j - v) h(u, v)$$

# Linear filtering

- Consider a 3x4 image and 2x2 kernel

1	5	1	1
1	2	8	2
0	0	1	8

Image

1	0
0	1

Kernel

3		

Output

Position 1:

1x1	5x0
1x0	2x1

# Linear filtering

- Consider a 3x4 image and 2x2 kernel

1	5	1	1
1	2	8	2
0	0	1	8

Image

1	0
0	1

Kernel

3	13	

Output

Position 2:

5x1	1x0
2x0	8x1

# Linear filtering

- Consider a 3x4 image and 2x2 kernel

1	5	1	1
1	2	8	2
0	0	1	8

Image

1	0
0	1

Kernel

3	13	3

Output

Position 3:

1x1	1x0
8x0	2x1



# Linear filtering

- Consider a 3x4 image and 2x2 kernel

1	5	1	1
1	2	8	2
0	0	1	8

Image

1	0
0	1

Kernel

3	13	3
1		

Output

# Linear filtering

- Consider a 3x4 image and 2x2 kernel

1	5	1	1
1	2	8	2
0	0	1	8

Image

1	0
0	1

Kernel

3	13	3
1	3	

Output

# Linear filtering

- Consider a 3x4 image and 2x2 kernel

1	5	1	1
1	2	8	2
0	0	1	8

Image

1	0
0	1

Kernel

3	13	3
1	3	16

Output

# Cross-correlation vs. convolution

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

$F[x,y]$

a	b	c
d	e	f
g	h	i

$H[u,v]$

		i	h	g		
		f	e	d		
		c	b	a		

$F[x,y] \otimes H[u,v]$

Cross-correlation: overlay filter on image

# Cross-correlation vs. convolution

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

$F[x,y]$

a	b	c
d	e	f
g	h	i

$H[u,v]$

		a	b	c		
		d	e	f		
		g	h	i		

$F[x,y] * H[u,v]$

Convolution: flip filter horizontally and vertically



# Summary

- Pixel operator: transform pixel based on its value
- Local operator: transform pixel based on its neighbours
- Convolution (and cross-correlation): operations that apply a linear filter to an image

# Common filters

# Practice with linear filters



Original

0	0	0
0	1	0
0	0	0


?

# Practice with linear filters




Original

0	0	0
0	1	0
0	0	0



Filtered  
(no change)

# Practice with linear filters



Original

0	0	0
0	0	1
0	0	0

?



# Practice with linear filters




Original

0	0	0
0	0	1
0	0	0



Shifted left  
By 1 pixel

# Practice with linear filters



Original

0	0	0
0	2	0
0	0	0

—


$\frac{1}{9}$

1	1	1
1	1	1
1	1	1

?

(Note that filter sums to 1)

# Practice with linear filters



Original

0	0	0
0	2	0
0	0	0

—

$\frac{1}{9}$


1	1	1
1	1	1
1	1	1




## Sharpening filter

- Accentuates differences with local average

# Common filters: Gaussian




$$G_{\sigma} = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$




Gaussian kernel  
Kernel size: 5 x 5 px  
 $\sigma = 1$

# Blur filters



Gaussian kernel




Average kernel






# Common filters: Sobel





# Common filters: Sobel



# What about colour?


- Consider a  $3 \times 4 \times 3$  image and  $2 \times 2 \times 3$  kernel



Convolution output?

# Designing filters

- How could you detect diagonal edges?





# Designing filters

- How could you simulate (linear) motion blur?



# Common filters

- Average/blur filters: average pixel values, blur the image
- Sharpening filters: subtract pixel from surround, increase fine detail
- Edge filters: compute difference between pixels, detect oriented edges in image

# Filters in practice

# Properties of linear filters


- Commutative:  $f * h = h * f$ 
  - Theoretically, no difference between kernel and image
  - But most implementations do care about order
- Associative:  $(f * h1) * h2 = f * (h1 * h2)$ 
  - Usually one option is faster than the other – allows for more efficient implementations
- Distributive over addition
  - $f * (h1 + h2) = (f * h1) + (f * h2)$
- Multiplication cancels out
  - $kf * h = f * kh = k(f * h)$





# Efficient filtering


- Multiple filters: generally more efficient to combine 2D filters ( $h1 * h2 * h3...$ ) and filter image just once

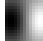



\*  =  
Gaussian  
blur filter



\*  =  
Horizontal  
derivative  
filter



\*  =  
Derivative-of-  
Gaussian filter



# Efficient filtering

- Separable filters: generally more efficient to filter with two 1D filters than one 2D filter
- For example, the 2D Gaussian can be expressed as a product of two 1D Gaussians (in x and y)

$$\begin{aligned} G_{\sigma}(x, y) &= \frac{1}{2\pi\sigma^2} \exp -\frac{x^2 + y^2}{2\sigma^2} \\ &= \left( \frac{1}{\sqrt{2\pi}\sigma} \exp -\frac{x^2}{2\sigma^2} \right) \left( \frac{1}{\sqrt{2\pi}\sigma} \exp -\frac{y^2}{2\sigma^2} \right) \end{aligned}$$

# Separable filters

2D convolution  
(center location only)

$$\begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} * \begin{bmatrix} 2 & 3 & 3 \\ 3 & 5 & 5 \\ 4 & 4 & 6 \end{bmatrix}$$

The filter factors into a product  
of 1D filters:

$$\begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} = \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} \times \begin{bmatrix} 1 & 2 & 1 \end{bmatrix}$$

Perform convolution along  
rows:

$$\begin{bmatrix} 1 & 2 & 1 \end{bmatrix} * \begin{bmatrix} 2 & 3 & 3 \\ 3 & 5 & 5 \\ 4 & 4 & 6 \end{bmatrix} = \begin{bmatrix} & 11 & \\ & 18 & \\ & 18 & \end{bmatrix}$$

Followed by convolution  
along the remaining column:

$$\begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} * \begin{bmatrix} & 11 & \\ & 18 & \\ & 18 & \end{bmatrix} = \begin{bmatrix} & & \\ & 65 & \\ & & \end{bmatrix}$$

# Convolution output size

- Valid convolution: the output image is smaller than the input image
- Why?

0	0	0	0	0	0	0
0	0	255	255	255	0	0
0	255	0	0	0	255	0
0	255	0	0	0	0	0
0	255	0	0	0	255	0
0	0	255	255	255	0	0
0	0	0	0	0	0	0

 $*$ 


$1/9$	$1/9$	$1/9$
$1/9$	$1/9$	$1/9$
$1/9$	$1/9$	$1/9$

 $=$ 

57	85	85	85	57
85	113	85	85	57
85	85	0	57	57
85	113	85	85	57
57	85	85	85	57

# Border handling

- Pad with constant value



Filter: Gaussian blur

# Border handling

- Wrap image



Filter: Gaussian blur





# Border handling

- Clamp / replicate the border value



Filter: Gaussian blur



# Border handling

- Reflect image



Filter: Gaussian blur





# Practical considerations

- Think about how to implement filters efficiently
  - Images are big, so efficient filtering can save a lot of time!
- Think about how to handle borders
  - No one-size-fits-all solution
  - Wrap is ideal for tiling textures (but not photos)
  - Clamp/replicate tends to work well for photos

# Summary

- Linear filters: first step of almost all computer vision systems
- Linear filters are just a first step – you can't build complex feature detectors from just linear filters