

Convolutional Neural Networks II

Semester 2, 2022

Kris Ehinger

Outline

- Downsampling
- Regularisation
- Training an image recognition CNN
- CNN results


Learning outcomes

- Implement max pooling and explain how/why downsampling is used in CNNs
- Implement regularisation methods and explain how/why they are used in CNNs
- Train a CNN and explain the design choices involved in the training process

Downsampling in CNNs

Convolutional neural network

“AlexNet”: Krizhevsky, Sutskever, & Hinton (2012)



Downsampling in CNNs

- It's common to downsample convolution layer output
- Reduces output size and number of computations needed in later layers
- Can also improve tolerance to translation – small changes in input won't change downsampled output

Strided convolution

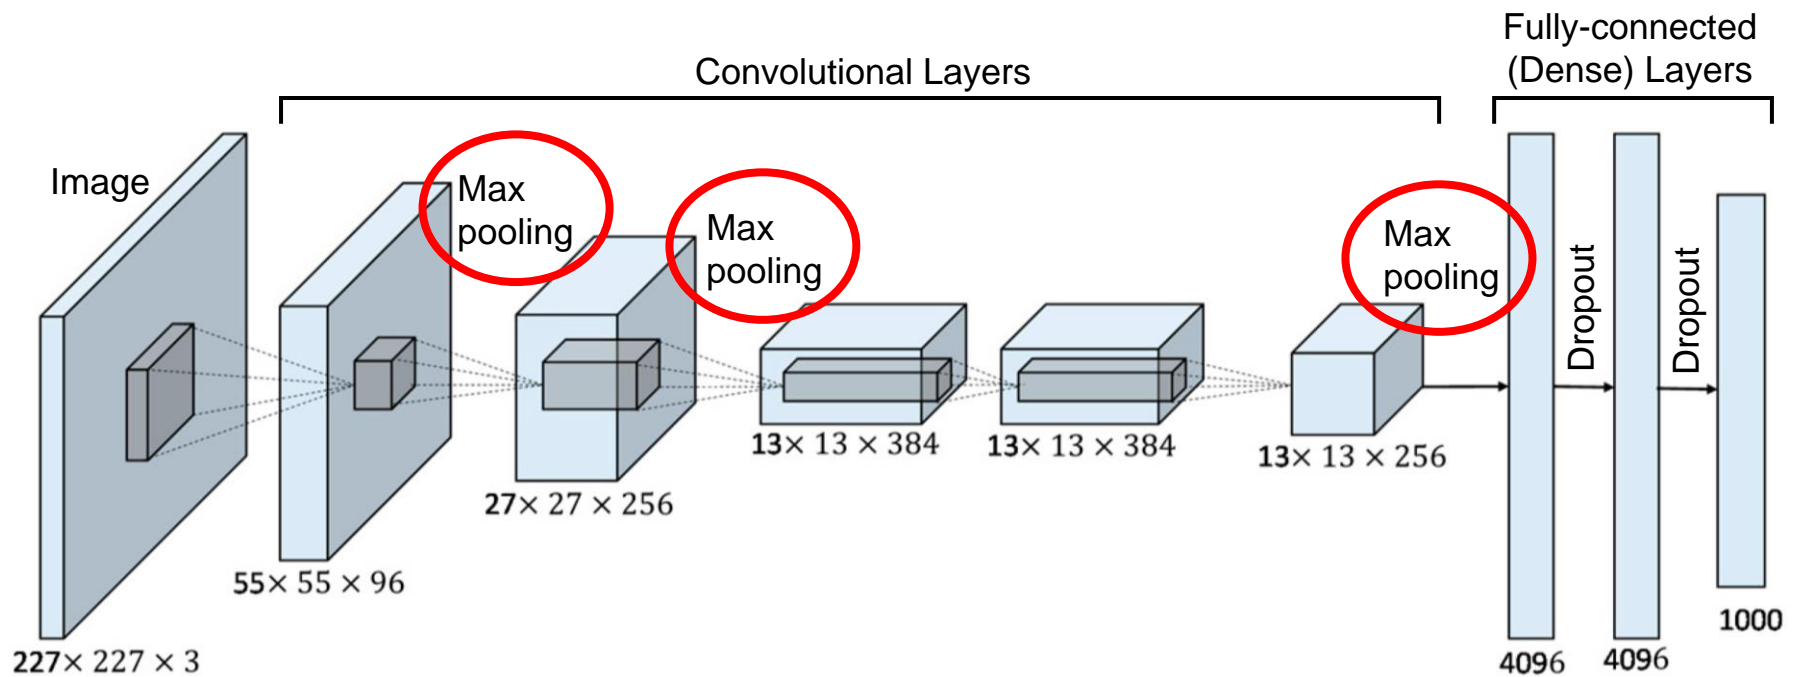
- Convolutional stride = distance between successive convolution windows
- In CNNs, stride can be >1
- Assuming no padding:
 - $\text{output_size} = \text{ceil}((\text{input_size} - \text{kernel_size} + 1)/\text{stride})$
- With padding:
 - $\text{output_size} = \text{ceil}(\text{input_size}/\text{stride})$

Strided convolution

- Advantage
 - Efficient – higher stride means fewer convolution operations
- Disadvantage
 - Kernel window skips over parts of the image, so important image features could be missed


Convolutional neural network

“AlexNet”: Krizhevsky, Sutskever, & Hinton (2012)




Max pooling

- After convolution, each activation map is separately downsampled
- Max pool stride determines the amount of downsampling ($\text{output_size} = \text{input_size} / \text{stride}$)




Max pooling

- Within a given window in the activation map, take the highest value and discard the rest



Average pooling

- Within a given window in the activation map, average the values



Max/average pooling

- Advantage
 - Max pooling is most likely to preserve the most important features, compared to strided convolution or average pooling
- Disadvantages
 - Average pooling “blurs” over features; important features may be lost
 - Pooling is slower than strided convolution

Summary

- Downsampling is common in CNNs to make computation more efficient in later layers
- Methods include strided convolution, max pooling, and average pooling

Regularisation in CNNs

Regularisation

- Due to the very high number of parameters, CNNs are prone to overfitting, even on large datasets
- Regularisation is usually needed to reduce overfitting
- Common options:
 - L1 or L2 regularisation
 - Dropout
 - Early stopping

L1, L2 regularisation

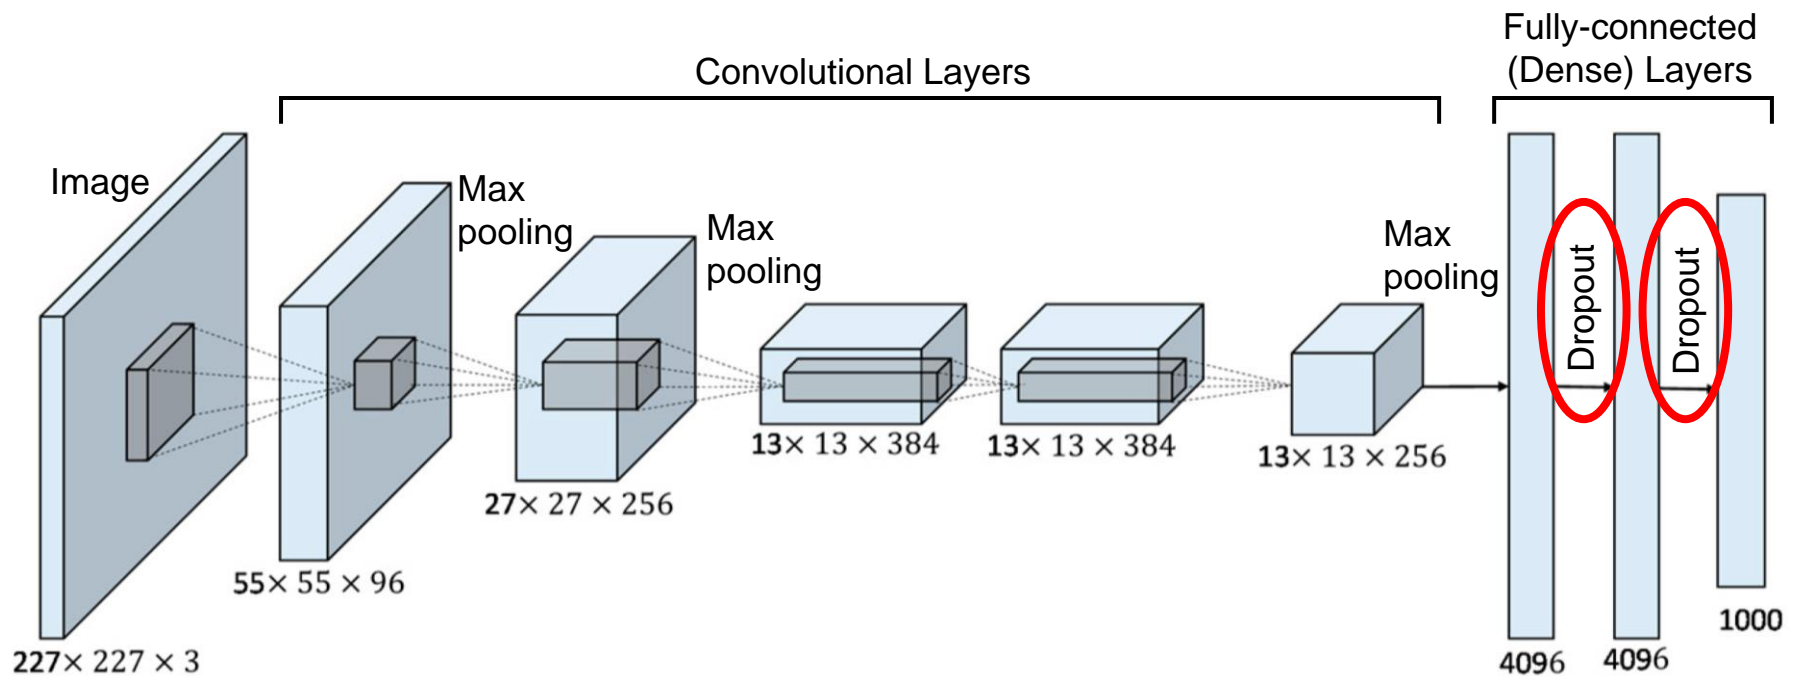
- Adds an additional term to the loss function that encourages smaller values for the network parameters
- L1 regularisation adds the term: $\sum_i |\theta_i|$
 - Penalise the sum of the absolute value of all parameters
 - Encourages sparse representation – many parameters should be 0
- L2 regularisation adds the term: $\sum_i \theta_i^2$
 - Penalise the sum of the squares of all parameters
 - Encourages small (but not 0) parameters

L1, L2 regularisation

- Free parameters when adding regularisation:
 - How much weight to give regularisation term vs. other terms in the loss function
 - Which layers to include in regularisation – all layers or just later layers?
 - Which parameters to include – sometimes only weights are included, not biases
- Adding regularisation tends to slow down training
- Too much regularisation can result in underfitting

Convolutional neural network

“AlexNet”: Krizhevsky, Sutskever, & Hinton (2012)



Dropout

- Randomly discard some neurons (set output = 0)
- Forces neurons to find useful features independently of each other
- Effectively, trains multiple architectures in parallel




Image: Srivastava, Hinton, Krizhevsky, Sutskever, Salakhutdinov (2014)

Dropout

- What percentage of neurons to drop is a free parameter (e.g., drop 50% or drop 20%)
- Can be applied to all layers, or just later layers
 - Different dropout percentages can be applied to different layers – typically later layers would have more dropout
- Adding dropout tends to slow down training
- Dropout is *only* used in training – when evaluating the network on new data (validation/test), all neurons are active

Early stopping

- Stop training the network when it shows signs of overfitting
- Monitor performance on a **validation** set
 - Subset of data not seen in training and not included in test set
 - During training, periodically check model's performance on the validation set – a decrease suggests overfitting
- Encourages smaller values for network parameters by keeping them close to their initial values (which are typically near 0)

Summary

- Regularisation is usually necessary to prevent overfitting
- Common options: L1 or L2 regularisation, dropout, and early stopping
- Frequently unclear which method (or combination) will work best for a given optimisation problem, so it's common to experiment and combine them

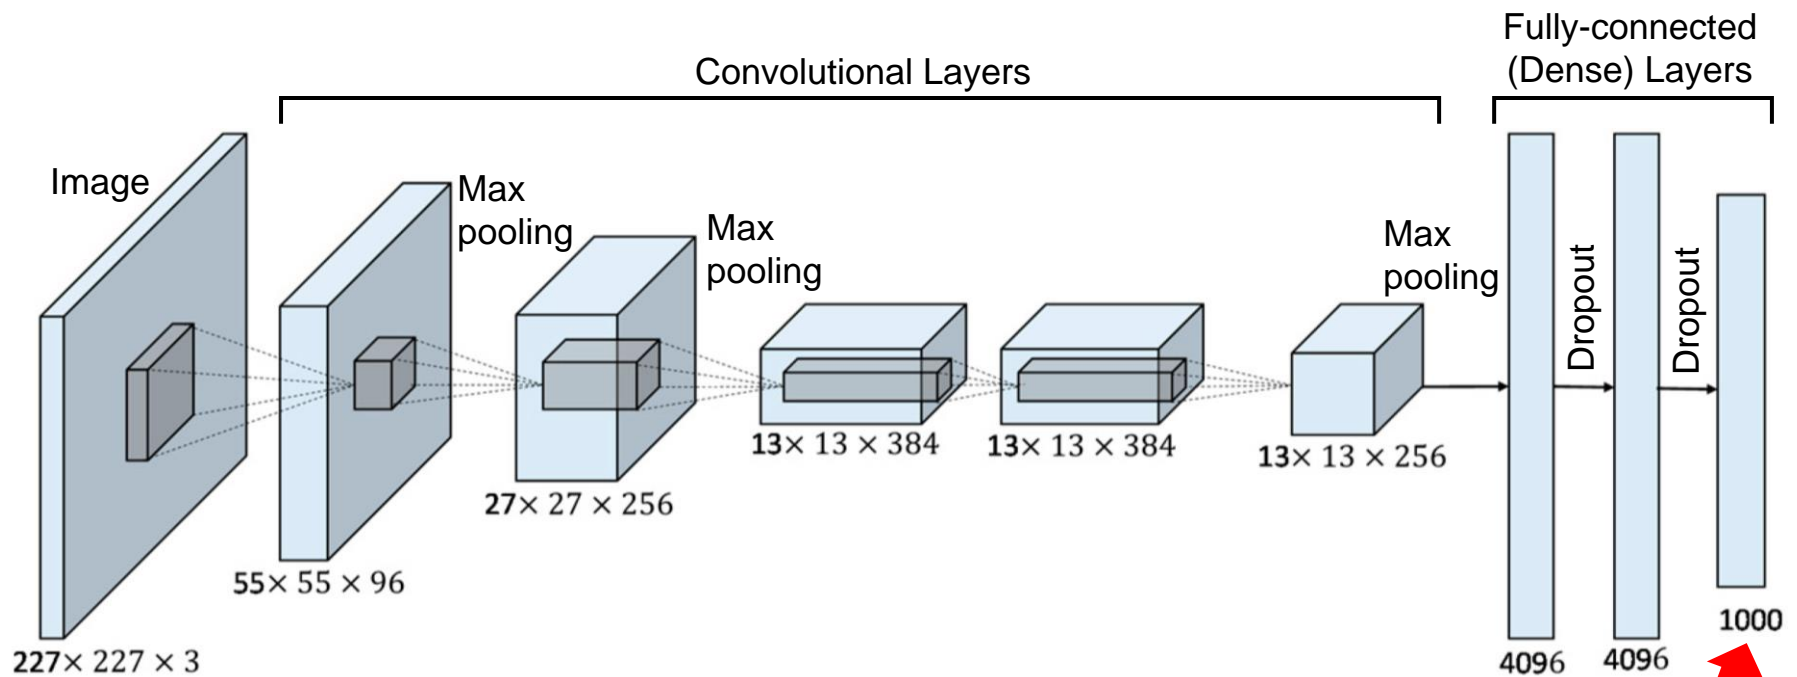
Training an Image Recognition CNN

CNN overview

- Typical architecture for image recognition:
 - Some number of convolutional layers, with downsampling
 - One or more fully-connected layers
 - Softmax output with cross-entropy loss
- Basic idea:
 - Do **feature embedding** in convolutional layers (transform images from pixels to useful high-level features)
 - Fully-connected layers are effectively a linear classifier (or MLP) to predict class from high-level features

Convolutional neural network

“AlexNet”: Krizhevsky, Sutskever, & Hinton (2012)



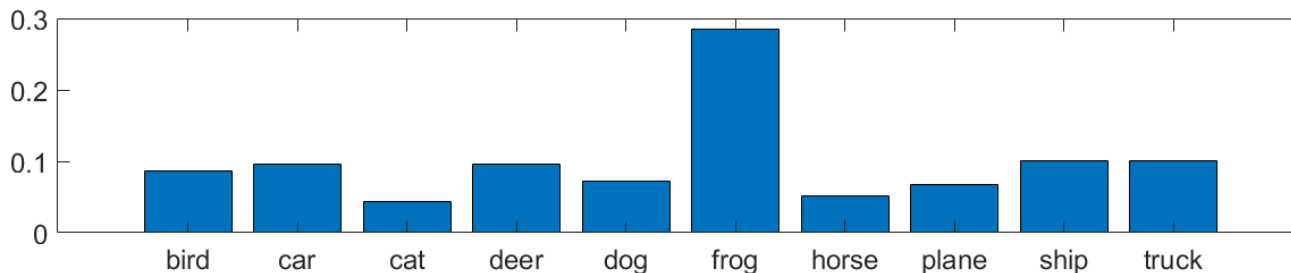
Output: softmax

Loss function: Softmax

- Apply softmax function to last layer's output:


$$\sigma(y_i) = \frac{e^{y_i}}{\sum_{j=1}^N e^{y_j}}$$

- Produces a vector that has the properties of a probability distribution:
 - All values in range 0-1
 - Values sum to 1



Loss function: Cross-entropy loss

- Measure of the difference between the model and ground truth probability distributions



- Cross-entropy loss between predicted class and ground-truth class:

$$E = -\frac{1}{N} \sum_{i=1}^N y_i \log(\hat{y}_i)$$

Annotations for the equation:


- N : N classes
- y_i : Ground truth probability (1 or 0)
- \hat{y}_i : Model probability from softmax

Training process

- Split data into train/validation/test sets
- Split training data into batches
- For $N = 1 - ?$
 - Preprocess a batch of image data
 - Classify batch, compute loss
 - Update model parameters with backprop
- Periodically check trained model's performance on the validation set (for early stopping)

Data preprocessing


- Image whitening – scale each image to 0-255 range, then normalise so each pixel has mean=0 and (optionally) std=1



```
X -= np.mean(axis=0, keepdims=True)
```

```
X /= np.std(axis=0, keepdims=True)
```

Data preprocessing



An input image (256x256)



Minus sign






The mean input image

A per-channel mean also works (one value per R,G,B).

Figure: Alex Krizhevsky

Data augmentation

- Manipulate training data to generate more samples
- Without data augmentation, even smaller networks (e.g., AlexNet) overfit to ImageNet



Data augmentation

- Common options:
 - Random crops (e.g., 224 x 224 from 256 x 256 images)
 - Horizontal reflection
 - Small colour/contrast adjustments (to simulate different camera settings or times of day)
- Less common:
 - Random rotation (e.g., +/- 15 degrees) ← slow
 - Random scale ← slow
 - Random occluders

Data augmentation

- Why not include other variations?
 - Vertical reflection
 - Large colour changes



Training process

- Initialise network weights and bias
 - Typically, weights initialised to small values from a Gaussian distribution around zero
 - Bias initialised to zero or small positive values
- Set training parameters
 - Batch size
 - Optimiser
 - Learning rate + decay
- Monitor training and validation loss
 - Stop training when validation loss no longer decreases

Batch size

- Batch size (or mini-batch size) = portion of the training data used to compute gradient for parameter update
- It's not computationally feasible to use the whole dataset to compute each update
- Dataset is randomly split into N batches of size b
- N updates = 1 epoch (every image has been seen once)

Batch size

- Smaller batch size
 - More updates (but these are faster to compute)
 - Noisier updates (high variance in gradient)
- Larger batch size
 - Fewer updates (but each update takes longer to compute)
 - More stable updates
- In practice, batch size tends to be limited by memory constraints

Optimiser

- Stochastic Gradient Descent (SGD)
- Root Mean Square Propagation (Rmsprop)
- Adaptive moment estimation (Adam)
 - Keep a moving average of the squared gradient/gradient to divide the learning rate
 - Different from SGD that maintains a single learning rate for different gradients with different magnitudes

Learning rate + decay

- Learning rate = how much to change network parameters on each update
 - Too high rate – unstable training
 - Too low rate – very slow learning




Image: <http://cs231n.stanford.edu/>

How long to train?

- Generally, train until models' performance on a validation set stops improving

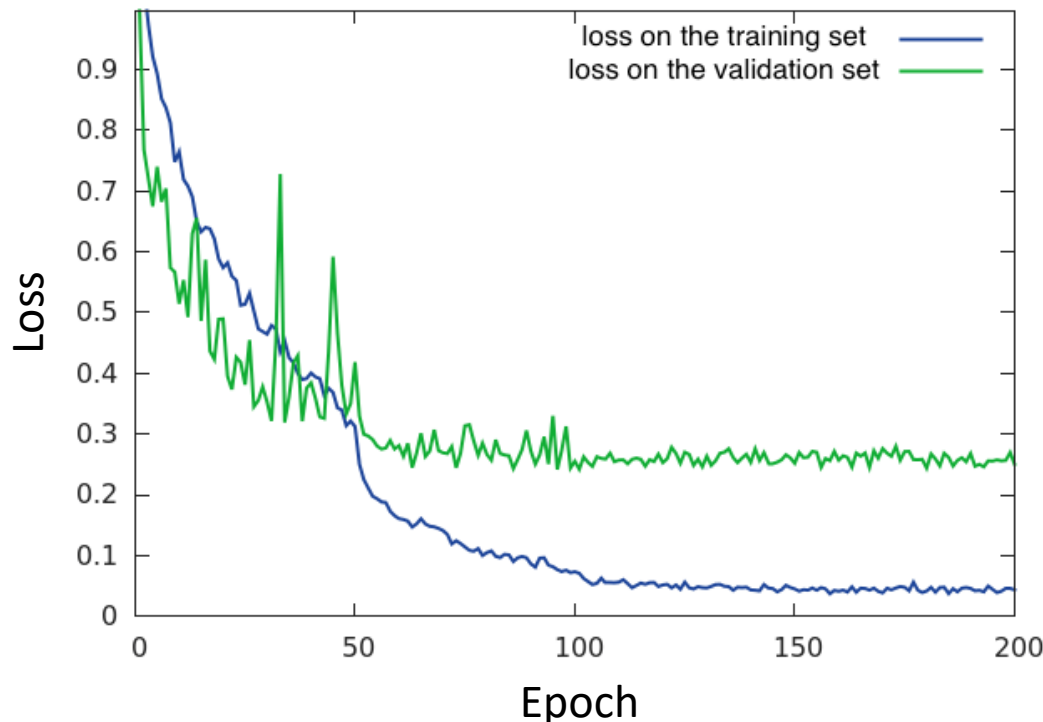


Image: Korbar, et al. (2017)

Summary

- Training CNNs can be difficult – parameter space is extremely large
- Data augmentation is usually required to avoid overfitting
- Hyperparameters (batch size, optimizer, learning rate) can affect how well the network learns

CNN results

Classification performance

ImageNet classification

Russakovsky et al. (2014):

“With a sufficient amount of training, a human annotator is still able to outperform the GoogLeNet result... by approximately 1.7%.”

2015: A MILESTONE YEAR IN COMPUTER SCIENCE





Image: <https://blogs.nvidia.com/wp-content/uploads/2016/01/2-milestone-web1.gif>

Classification performance



<https://paperswithcode.com/sota/image-classification-on-imagenet>

Classification performance

Easiest
classes:



Hardest
classes:



Classification errors

ImageNet Classification Failures (GoogLeNet 2014)



ruler

pencil box
rubber eraser
ballpoint pen



king crab

pizza
strawberry
orange



sidewinder

maze
gar
valley



saltshaker

pill bottle
water bottle
lotion



reel

stethoscope
whistle
ice lolly




hatchet

vase
pitcher
coffeepot

Classification errors

- Performance of top models on ImageNet vs. ImageNetV2
- Drop of about 10% suggests some overfitting to ImageNet



Summary

- CNNs are the state-of-the-art for image classification, exceeding human performance on ImageNet
- CNN classification errors are often understandable (odd views, small objects), which suggests they learn reasonable features for this task
- But they do show some odd failures, like poor generalisation to ImageNetV2