

Texture

Semester 2, 2021

Kris Ehinger

Shape and texture



Inpainting demo

- <https://www.nvidia.com/research/inpainting/index.html>

Outline


- What is texture?
- Texture synthesis
- Texture transfer

Learning outcomes

- Explain parametric and non-parametric methods for representing and synthesizing texture
- Explain common applications of texture synthesis
- Explain an algorithm for texture transfer (Neural Style Transfer)

What is texture?


What is texture?



A



B



C

What is texture?

- A definition from image processing:



Texture is an region with **spatial stationarity** (same statistical properties everywhere in the region)

- A definition from computer graphics:


Texture is a 2D surface applied to a 3D model

Types of texture

- Periodic texture – has a subregion that repeats in a regular pattern







- Stochastic (aperiodic) texture – generated by a random process



Images: Describable Textures Dataset (DTD), Efros & Freeman (2001)

Texture models

- **Parametric** models:
represent texture with a
set of adjustable
parameters
- **Non-parametric**
(stitching) models:
represent texture as
image patches




Why model texture?

- Texture synthesis – create more of a texture
 - Textures for computer graphics, video games, etc.
 - Image inpainting
- Texture transfer
 - Artistic effects
 - Online shopping


Texture synthesis

Image stitching approach



How do you fill in the missing data?
Look at a neighbourhood around this patch...

Image stitching approach



How do you fill in the missing data?


Look at a neighbourhood around this patch...

and find similar neighbourhoods in other parts of the texture.



Image: Ehinger & Rosenholtz (2012)

Image stitching approach



Original patch:




Similar patches:



Select probable value for the missing pixel, based on the similar patches:



Non-parametric texture synthesis



1. Randomly sample a small (e.g., 3 x 3 pixel) patch from the original image
 2. Spiral outward, filling in missing pixels by finding similar neighborhoods in the original texture
- (Neighbourhood size is a free parameter that specifies how stochastic the texture is)

Neighbourhood size




Image quilting

- Efficient patch-based texture synthesis (Efros & Freeman, 2001)
- Use existing patches of texture to synthesis more texture; main problem is connecting them together without visible artefacts/seams
- “Corrupt Professor’s Algorithm”
 - Plagiarize as much of the source image as you can
 - Then try to cover up the evidence

Image quilting algorithm

- Choose patch and overlap size
- Initialize with a random patch
- For each subsequent patch:
 - Find a patch in the original texture that is most similar to this region, considering only the pixels in the overlap region
 - Seamlessly paste in patch by cutting along a path with minimum overlap error




Image quilting algorithm





Image quilting

Overlapping blocks



Vertical boundary





Diagram illustrating the calculation of overlap error. Two overlapping blocks are shown, with a subtraction operation ($-$) applied to the overlapping region. The result is squared (2) and equated ($=$) to a vertical strip showing the error, with a red line indicating the boundary.

$$\left[\text{Block 1} - \text{Block 2} \right]^2 = \text{Overlap error}$$

Min. error boundary



Graph cuts

- Represent neighbouring pixels as a graph
- Edge weight = overlap error
- Problem: Find path through graph with minimum total overlap error






Image quilting results



Original

Synthesis



Original

Synthesis




Image quilting results

The key to solving many algorithmic problems is to think of them in terms of graphs. Graph theory provides a language for talking about the properties of relationships, and it is amazing how often messy applied problems have a simple description and solution in terms of classical graph properties.

Designing truly novel graph algorithms is a very difficult task. The key to using graph algorithms effectively in applications lies in correctly modeling your problem so you can take advantage of existing algorithms. Becoming familiar with many different algorithmic graph *problems* is more important than understanding the details of particular graph algorithms, particularly since Part II of this book will point you to an implementation as soon as you know the name of your problem.

Here we present basic data structures and traversal operations for graphs, which will enable you to cobble together solutions for basic graph problems. Chapter 6 will present more advanced graph algorithms that find minimum spanning trees, shortest paths, and network flows, but we stress the primary importance of correctly modeling your problem. Time spent browsing through the catalog now will leave you better informed of your options when a real job arises.

ing s, particularly since Part II of with manyaph algorithms, particularly since Part II of this
book on as you know the name of yanding thenentation as soon as you know the name of your p
oblemes and traversal operations foris book wasic data structures and traversal operations for grap
s, whiutions for basic graph proble problem.bble together solutions for basic graph problems. C
naptgorithms that find minimum saphs, whicvanced graph algorithms that find minimum span
ur pran take advantage of existing algorithms. Becoming familiar with mans for basic graph p
with algorithmic graph *problems* is more important than understanding thms that find minim
andif particular graph algorithms, particularly since Part II of this book wiess the primary imp
s boot to an implementation as soon as you know the name of your problem. ing through the ca
probwe present basic data structures and traversal operations for graphs, whiof this bocniliar with
aphs.ble you to cobble together solutions for basic graph problems. Chapteryour proUnderstandi
Chapnt more advanced graph algorithms that find minimum spanning treeor graphs, of this bo
lving paths, and network flows, but we stress the primary important tof correctems. Chapyour pro
theory providesions lies ins lies in correctly mophs, whichrly since Part II of spanningfor graphs
it is amazing lg algorithmalgorithms. Becoming Chapter know the name of yance of conlems. Cha
lution in terms is more imore important thaning treeaversal operationsfor, now will spanning
y novel graph algorithms is a very difficult task. The key r basic graph problems. Chapteble toge
effectively in applications lies in correctly modeling your that find minimum spanning trvanced gr
dvanantage of existing algorithms. Becoming familiar wi the primary importance of correctwork flo
nic graph *problems* is more important than understan through the catalog now will lem. Time
ar graph algorithms, particularly since Part II of this bs.
plementation as soon as you know the name of your prask. The key to usingey to usiney to usin
t basic data structures and traversal operations for graphsodeling your problemur problemr problem
cobble together solutions for basic graph problems. Clg familiar with mar with manywith man
advanced graph algorithms that find minimum spanningn understanding thstanding thstanding th
l network flows, but we stress the primary importance ofrt II of this book wills book wil book wi

Image quilting results





Image inpainting

- Similar idea to fill in missing regions of an image:
 - Find a similar patch in *another* image
 - Paste in patch with an error-minimizing cut



Parametric texture synthesis



- Alternative to stitching approaches: represent texture with a number of parameters
- To synthesize texture, coerce a noise image to match the required parameters (usually through gradient descent)
- What parameters are needed to define a texture?

Fourier magnitude?





Image

Week 8, Lecture 1



Magnitude + random phase

COMP90086 Computer Vision





Phase + random magnitude



30

Fourier texture synthesis

- Synthesize texture by matching Fourier magnitude
- Okay results for some simple textures, but doesn't work well in general



Original image and power spectrum





Synthesized image and power spectrum

Images: Portilla, Navarro, Nestares, & Tabernero (1996)

Colour and edges?

- Textures could be defined as a distribution over simple features, like colour and edge orientation at various scales
- Synthesize texture by matching the distribution



Edge detector response at various orientations/scales

Distribution-matching results

Original



Synthesis



Original



Synthesis




Images: Heeger & Bergen (1995)

More complex statistics?

- Simple distributions of features are not sufficient
- Also need to represent feature co-occurrence



Image




Matching feature distributions
(Heeger & Bergen, 1996)



Matching feature distributions
and correlations
(Portilla & Simoncelli, 2000)

Texture synthesis results




Images: Portilla & Simoncelli (2000)

Even more complex statistics?

- The set of statistics needed to represent real images may be very complex
- Instead of modelling statistics by hand, represent texture as the feature response in the layers of a neural network trained on ImageNet classification

conv1




conv2




conv4



conv5

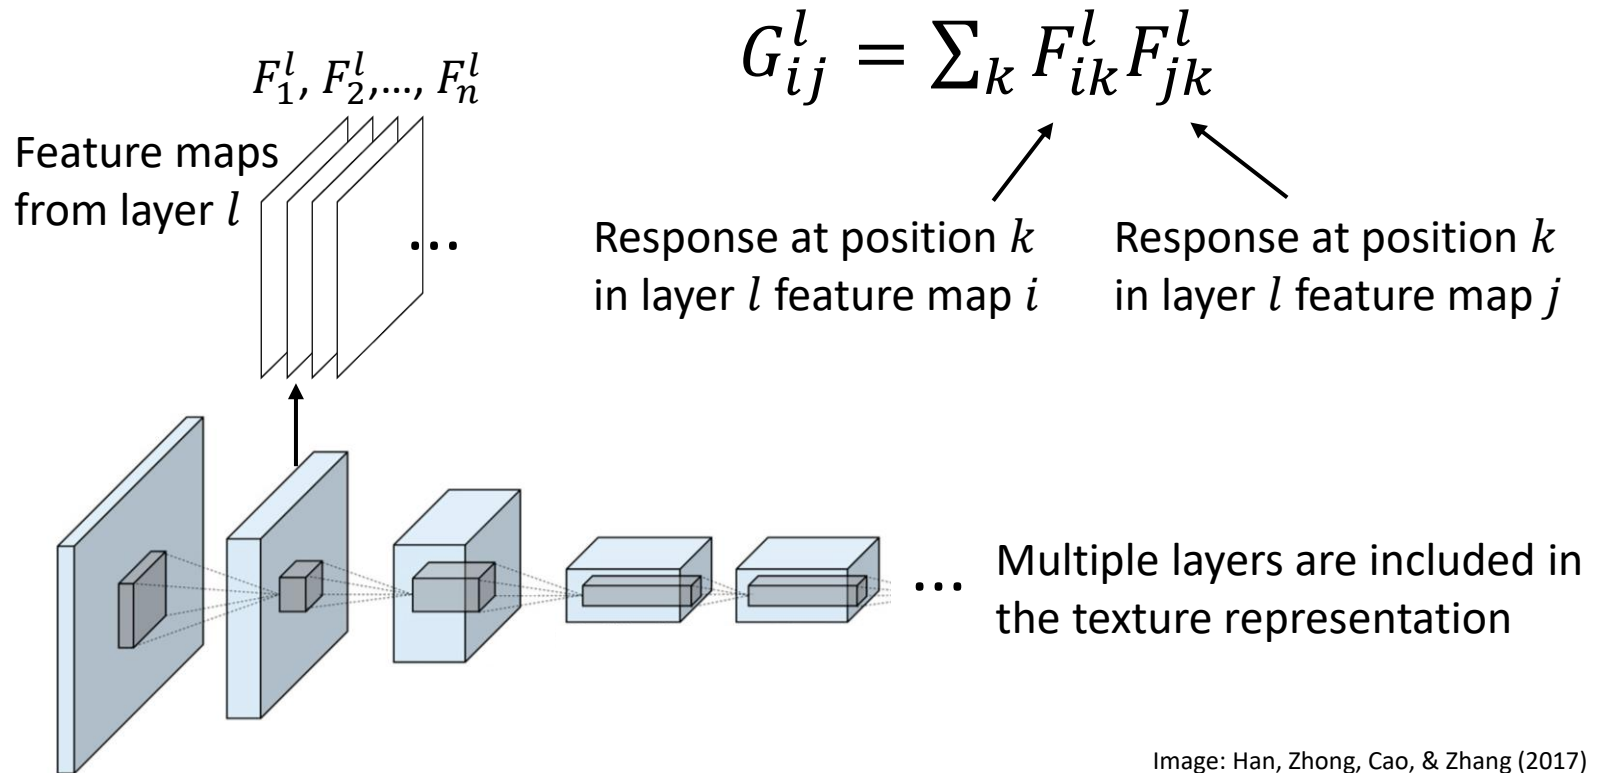


original

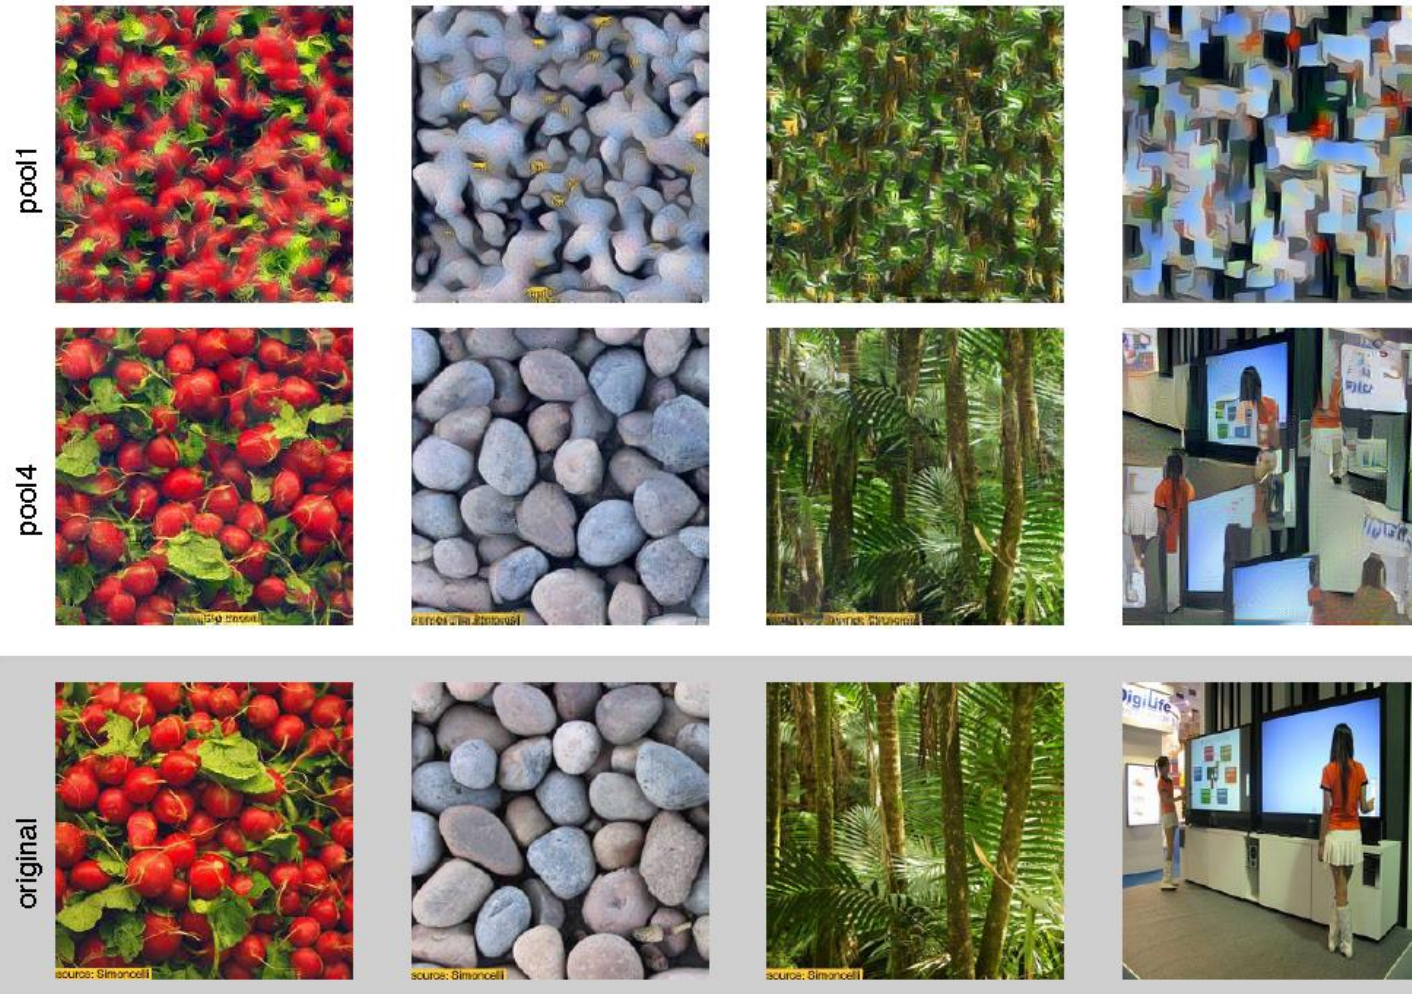


Feature correlations

- Texture is represented as the correlations between feature maps at a layer of the neural network:



Texture synthesis results



Texture synthesis results

Synthesis



Original image



Images: <http://bethgelab.org/deeptextures/>

Summary

- Non-parametric texture synthesis is based on copying texture patches
 - Works very well on periodic textures
 - Disadvantage: No model of texture parameters
- Parametric texture synthesis represents textures in terms of a set of parameters
 - Most methods work better on stochastic textures
 - Disadvantage: Even very complex models (e.g., based on neural networks) may be incomplete

Texture transfer

Texture transfer

- Render an image in the style of another image:



Content image

+



Style image



Image: Gatys, Ecker, & Bethge (2016)

Neural style transfer algorithm

- Both images (content, style) are run through a VGG network trained on ImageNet
- Content is represented as the responses from a layer of the neural network
- Style is represented as the correlations between feature maps at a layer of the neural network
- Use gradient descent to find an image that matches both content and style

Neural style transfer

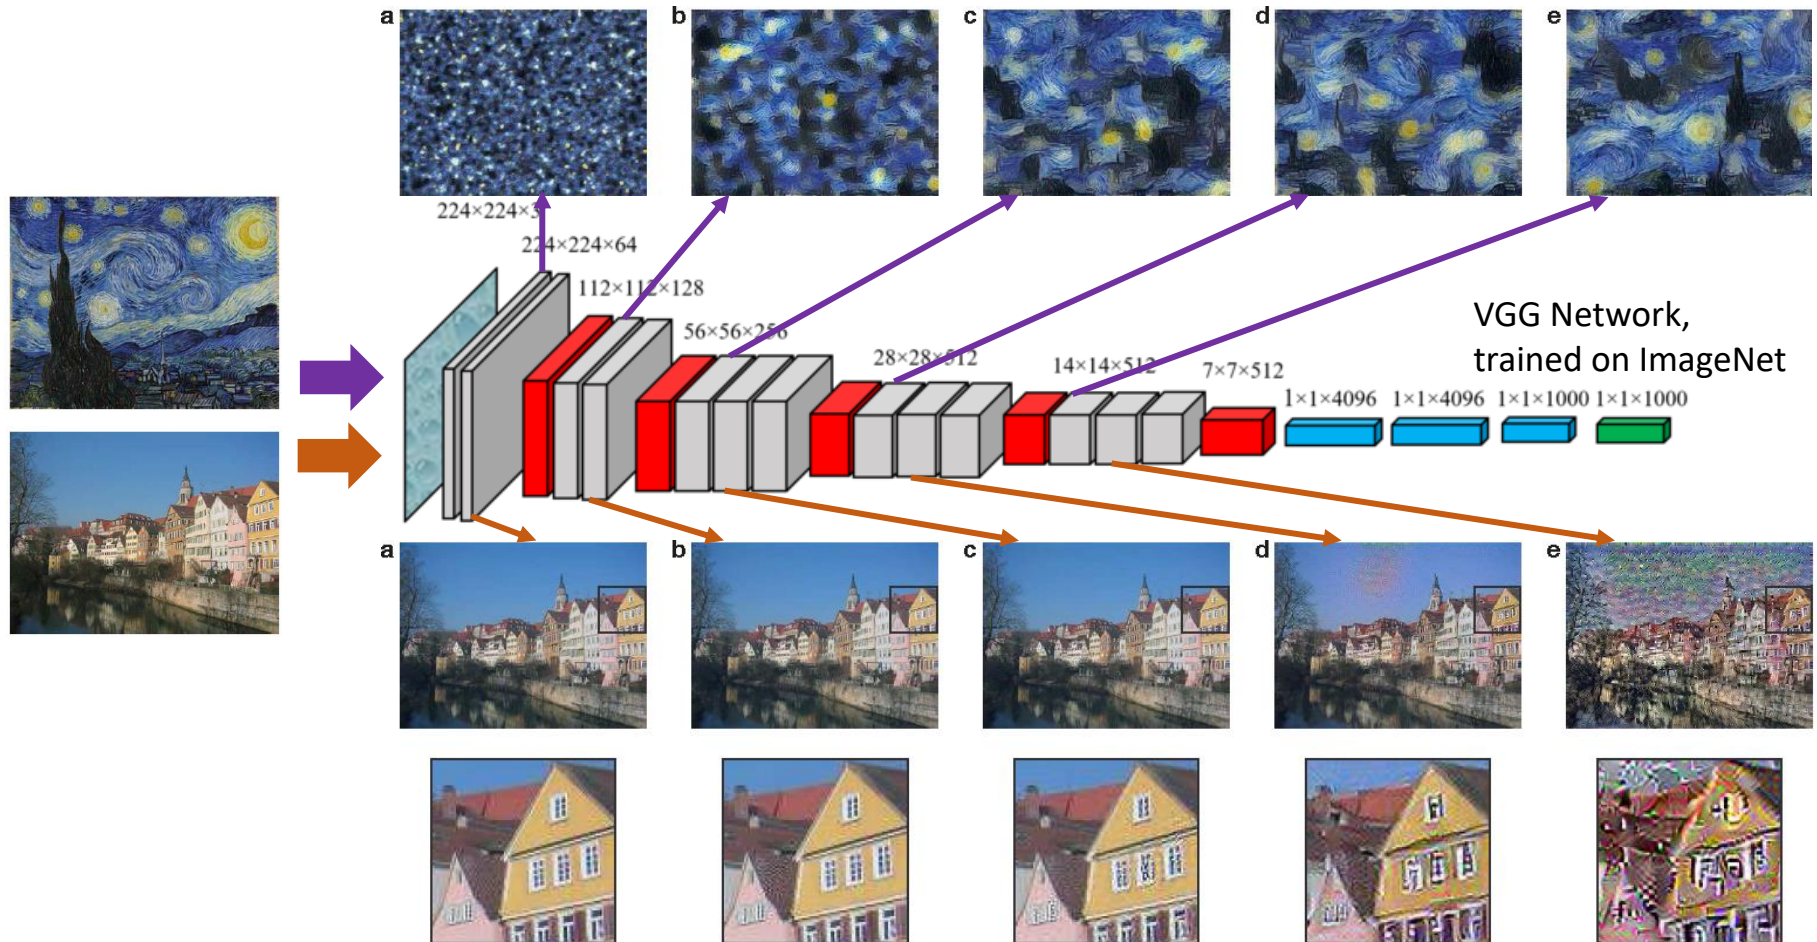


Image: Gatys, Ecker, & Bethge (2016)

Style transfer parameters

- Loss is sum of loss from content reconstruction and style reconstruction
- Relative weight of content vs. style is a free parameter:



Style image



← More weight on style

More weight on content →

Image: Gatys, Ecker, & Bethge (2016)

Style transfer parameters


- Content and style can be matched at any combination of layers
- Generally, match content at higher layers, and style across all layers



Style image



Original



Matched at layer conv2_2



Matched at layer conv4_2

Image: Gatys, Ecker, & Bethge (2016)

Texture transfer results




Image: Gatys, Ecker, & Bethge (2016)

Summary

- Texture transfer – render an image in the style of another image
- Image content represented by neural network responses
- Texture represented by correlations of feature maps across multiple layers of a neural network

Summary

- Texture can be defined in different ways, but generally captures 2D/surface aspects of an image
- Texture representations are useful for texture synthesis and texture transfer
- Applications:
 - Image inpainting
 - Computer graphics
 - Art