

Image Segmentation I

Semester 2, 2021

Kris Ehinger


Demo

- <https://chocopoule.github.io/grabcutweb/>



Segmentation



Separate image
into different
regions (objects,
textures)



Semantic segmentation



Separate image
into different
labelled regions

Instance segmentation

Semantic
segmentation



Instance
segmentation







Image segmentation



Input: Image



Clustering?
Graph cuts?
Classification?



Output: Pixel classification
(and, optionally, labels)

Outline


- Pixel clustering
- Superpixel segmentation
- Graph-based segmentation

Learning objectives

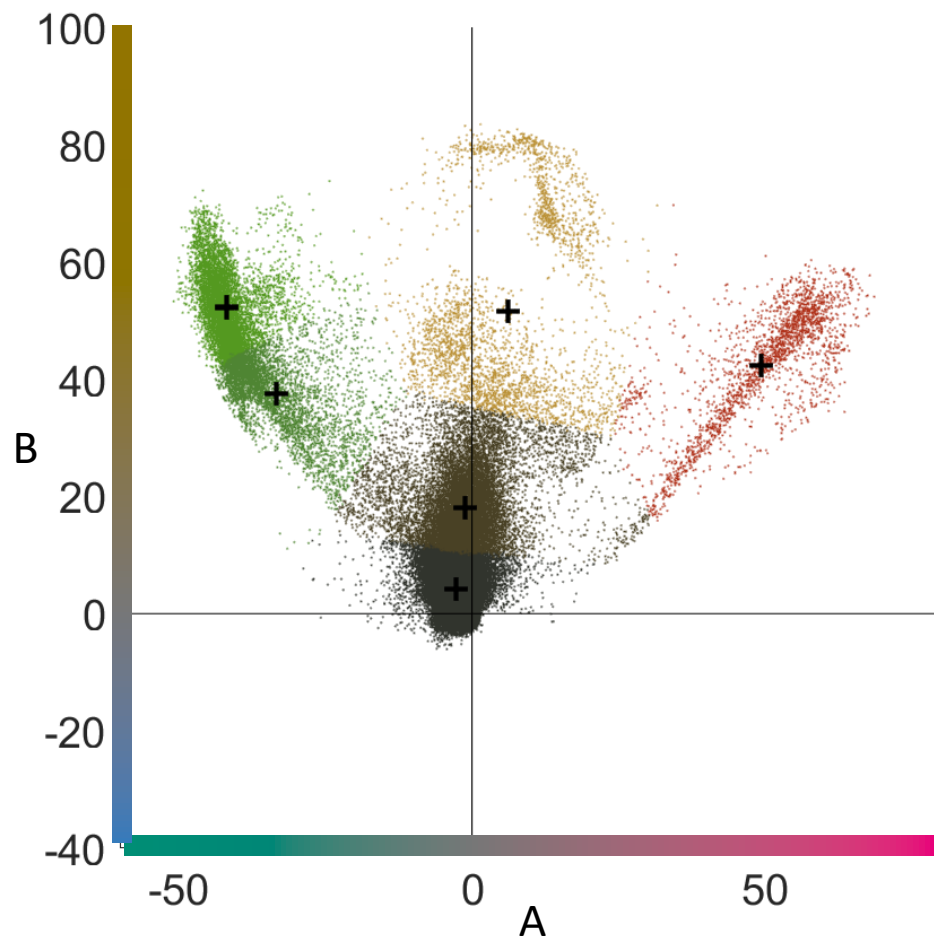
- Implement clustering algorithms for segmentation and compare/contrast clustering methods
- Implement an algorithm for computing superpixels and explain their common applications
- Explain graph-based methods for image segmentation

Pixel clustering


Colour clustering




K-means ($k = 6$)






Gaussian mixture model ($k = 6$)



Mean shift (bandwidth = 7)

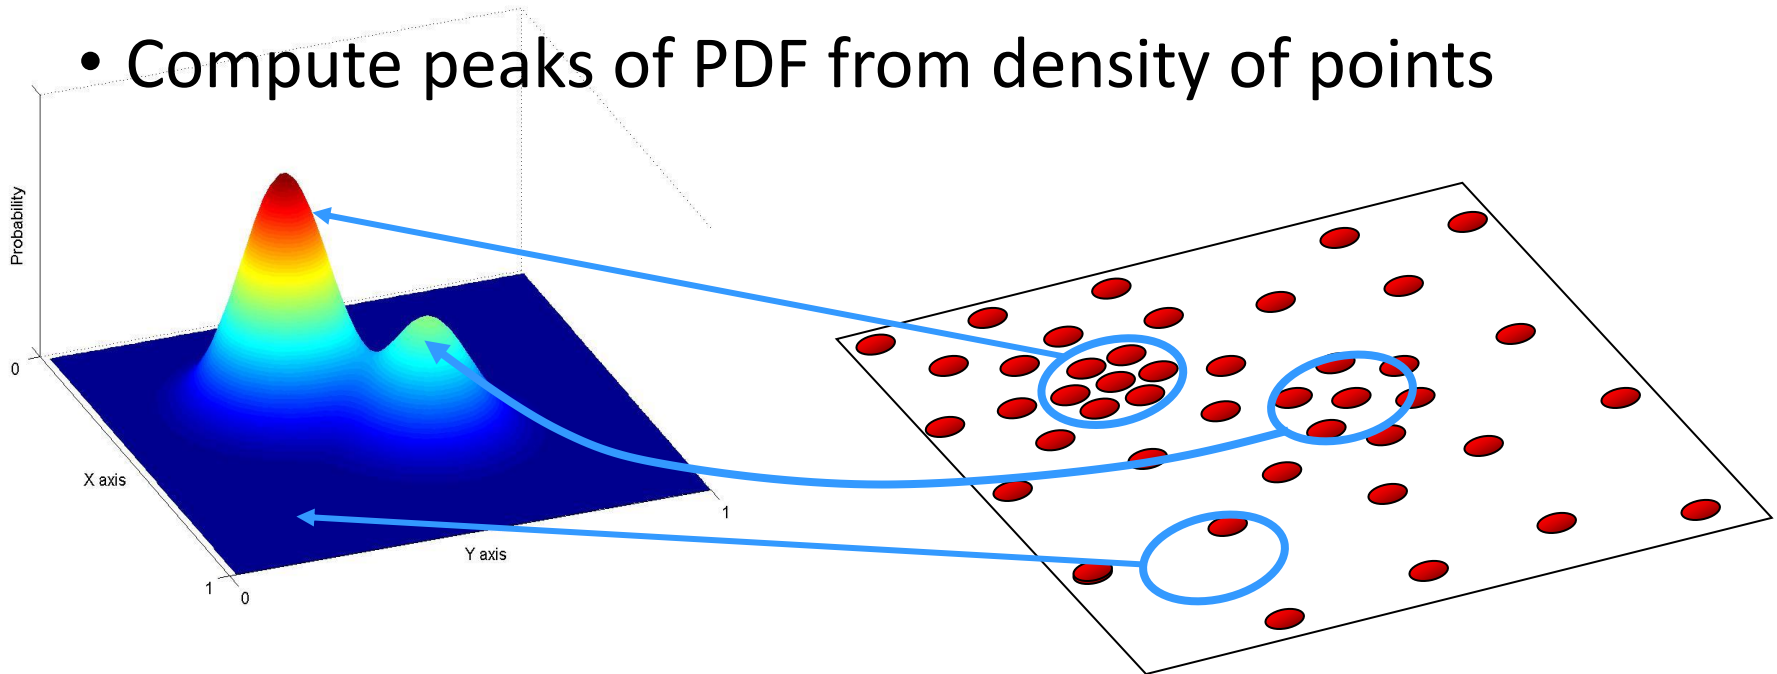


Mean shift clustering



Mean shift clustering

- Assume points are samples from an underlying probability density function (PDF)
- Compute peaks of PDF from density of points




Assumed Underlying PDF


Real Data Samples

Image: Y. Ukrainitz & B. Sarel


Mean shift




Mean shift




Mean shift




Mean shift




Mean shift



Mean shift



Mean shift




Mean shift algorithm

- Compute mean shift vector $\mathbf{m}(\mathbf{x})$
- Translate kernel window by $\mathbf{m}(\mathbf{x})$

Gaussian kernel:

$$g(\mathbf{x}) = \frac{1}{\sqrt{2\pi}} e^{-\frac{\mathbf{x}^2}{2}}$$

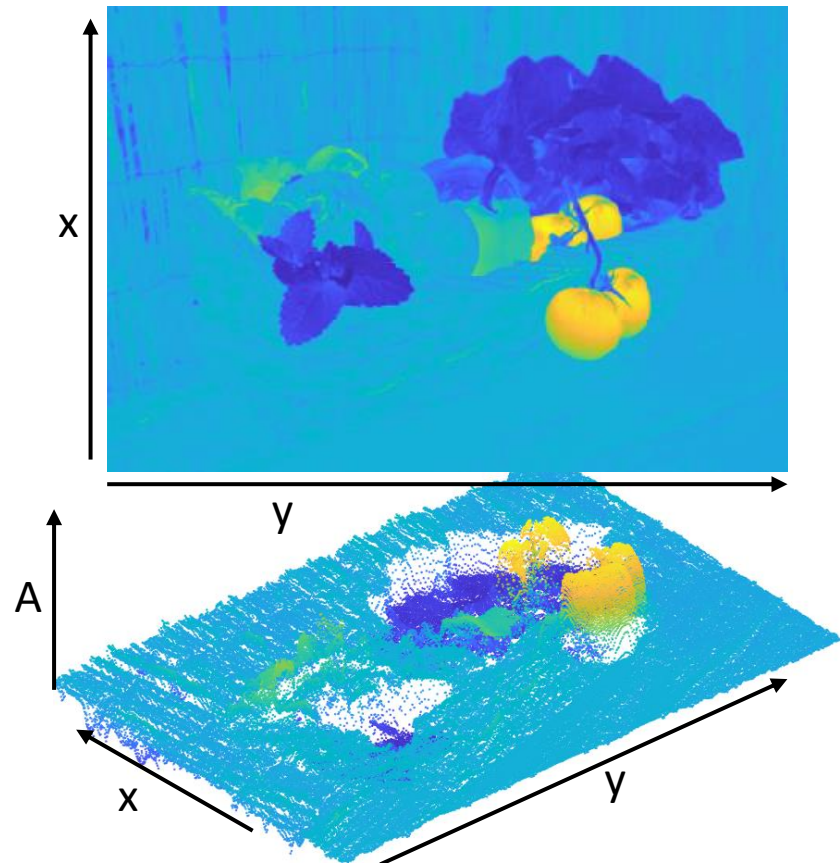


Mean shift algorithm

- For each point:
 - Centre a window on that point
 - Compute the mean of the data in the search window
 - Centre the search window at the new mean location
 - Repeat (b,c) until convergence
- Assign points that lead to nearby modes to the same cluster
- Free parameters: kernel (commonly Gaussian), bandwidth

Mean shift segmentation

- Cluster in spatial+colour space; e.g.: (x, y, R, G, B) or (x, y, L, A, B) coordinates



Mean shift parameters



Increasing bandwidth →

Increasing spatial bandwidth relative to colour →

Summary

- Pixel clustering is a fast, simple approach to image segmentation
- Example: mean shift clustering in the colour+spatial domain
 - Automatically discover number of clusters; no need to choose k
 - But do need to choose bandwidth
- Pixel clustering separates colour regions – regions may not correspond to objects

Superpixels

Superpixels

- **Oversegmentation** methods segment image into regions that are smaller than objects
 - Objects are separated from background
 - But objects are also separated into many parts
- Superpixels = groups of adjacent pixels with similar characteristics (e.g., colour)

Superspixel segmentation




Image: <https://www.epfl.ch/labs/ivrl/research/slic-superpixels/>

SLIC superpixel algorithm

- Initialise cluster centres on non-edge pixels:
 - Initialise k cluster centres $c_k = [x_k, y_k, l_k, a_k, b_k]$ by sampling the image in a regular grid
 - For each centre c_k , check an $N \times N$ neighbourhood around c_k to find the pixel with lowest gradient. Set c_k to this pixel's $[x, y, l, a, b]$.

SLIC superpixel algorithm

- For each cluster centre c_k :
 - In a $2M \times 2M$ square neighbourhood around c_k , measure pixel similarity to c_k
 - Assign pixels with similarity $<$ threshold to cluster k
 - Compute new cluster centre c_k
- Repeat until average change in cluster centres (L1 distance) falls below a threshold

• Similarity measure: $D = D_{lab} + \frac{\alpha}{M} D_{xy}$

$$D_{lab} = \sqrt{(l - l_k)^2 + (a - a_k)^2 + (b - b_k)^2}$$
$$D_{xy} = \sqrt{(x - x_k)^2 + (y - y_k)^2}$$


α = weighting parameter

SLIC superpixel algorithm

- Similarity metric does not guarantee that clusters will be connected pixels
- To enforce connectivity, pixels not connected to main cluster are re-assigned to closest adjacent cluster

Supapixel methods

Graph-based methods




Felzenszwalb &
Huttenlocher
(2004)

Veksler, Boykov,
& Mehrani
(2010) –
spatially
compact

Veksler, Boykov,
& Mehrani
(2010) –
constant colour

Gradient-descent-based



QuickShift
(Vedaldi &
Soatto, 2008)

SLIC

Supapixel applications



- Supapixels are a multipurpose intermediate image representation
- More compact representation for algorithms with high time complexity (600x800 pixels -> 200 supapixels)
- Common application: object segmentation
 - Oversegment image
 - Combine supapixels to find objects

Supapixel merging

- Region Adjacency Graph (RAG)
 - Vertices = image regions (pixels or superpixels)
 - Edge weights = difference between regions
- To merge superpixels:
 - Identify edges below a threshold and re-label superpixels connected by these edges as one region
 - Or iteratively:
 - Find lowest-weight edge, relabel connected superpixels as one region
 - Recompute RAG, repeat until a criterion is met (e.g., all edges above a threshold)



Superspixel merging

Superspixels



Find lowest-weight edge, merge superspixels

RAG



Recompute RAG, repeat


Summary

- Superpixels = regions of similar pixels, produced through oversegmentation
- Various algorithms for computing superpixels, SLIC is one common option
- Superpixels are a compact, intermediate representation used as a first step for:
 - Segmentation (especially graph-based methods)
 - Object detection/localisation
 - Video tracking

Graph-based segmentation

Images as graphs

- Represent image as a graph $\mathbf{G} = (\mathbf{V}, \mathbf{E})$
 - Vertices = image regions (pixels or superpixels)
 - Edge weights = similarity between regions



$$E_{12} < E_{23}$$

Graph cuts


- Consider image as a fully-connected graph
- Partition graph into disjoint sets A,B to maximize total edge weight = remove low-weight edges between dissimilar regions
- Minimize value of cut:

$$cut(A, B) = \sum_{u \in A, v \in B} w(u, v)$$

↖
Weight of edge
connecting u and v

Graph cuts

- Not ideal for image segmentation – tends to create small, isolated sets



Edge weight = $1/\text{distance}$

Normalised cuts

- Instead of minimizing cut value, minimize cut value as a fraction of total edge connections in entire graph (normalised cut)
- Normalised cut (Shi & Malik, 2000):

$$\begin{aligned} Ncut(A, B) &= \frac{cut(A, B)}{assoc(A, V)} + \frac{cut(A, B)}{assoc(B, V)} \\ &= \frac{\sum_{u \in A, v \in B} w(u, v)}{\sum_{u \in A, t \in V} w(u, t)} + \frac{\sum_{u \in A, v \in B} w(u, v)}{\sum_{v \in B, t \in V} w(v, t)} \end{aligned}$$

Normalized cuts results



Image: D. Hoiem

GrabCut

- Segments image pixels into just two classes: foreground (object) and background
- Uses colour clustering + graph cuts to find optimal classification of pixels into each class




Rother, Kolmogorov, & Blake (2004)

GrabCut algorithm

- Requires user to initialise algorithm with a bounding box



Outside box = background pixels



Inside box = treat
as foreground
pixels initially

GrabCut algorithm

- For each class (foreground, background), represent distribution of pixel colour as a Gaussian mixture model (GMM)
- Represent image pixels as a graph (8-way connectivity)



GrabCut algorithm

- Denote the pixel graph as \mathbf{G} and the GMM as θ
- α indicates label of each pixel (foreground or background)
- Iterate until convergence:

- Find graph cut (label assignment) to minimize

$$E(\alpha, \theta, \mathbf{G}) = U(\alpha, \theta, \mathbf{G}) + \gamma V(\alpha, \mathbf{G})$$

-log likelihood of cluster
assignments in GMM


Weighting
parameter

Smoothness penalty based on colour
similarity, applied to neighbouring
pixels with different labels in α


- Recompute GMM for new label assignment

GrabCut example


Initialisation



E over iterations



Initial GMM



Final GMM


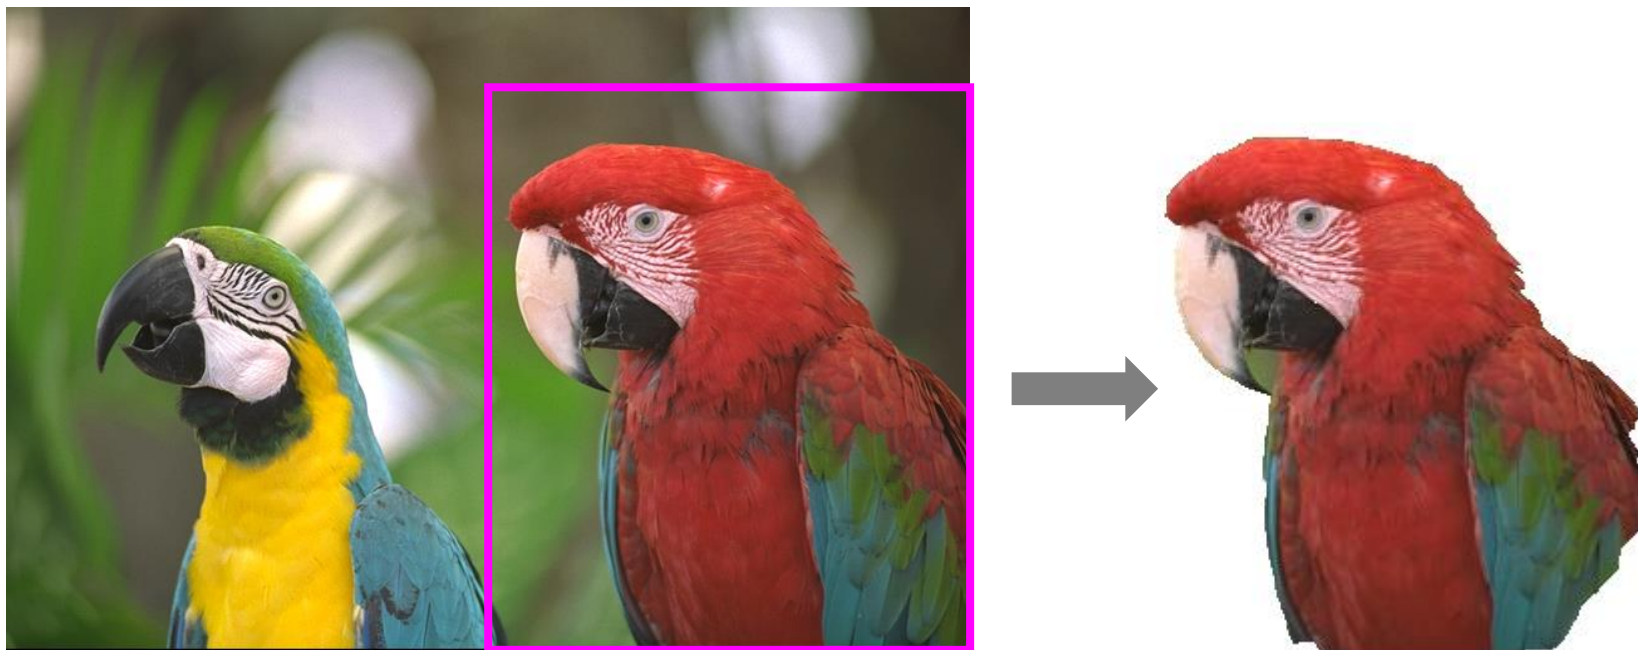


Image: Rother, Kolmogorov, & Blake (2004)

GrabCut result



Summary

- Graph-based methods represent an image as a graph (of pixels or superpixels)
- Segmentation removes edges to break graph into subgraphs, generally trying to optimize:
 - Similarity within connected region
 - Dissimilarity across disconnected regions
 - Smoothness/connectivity of connected regions
- Normalized cuts – segment into multiple regions
- GrabCut – segment into foreground/background

Summary

- Various ways to approach image segmentation, but many methods use some combination of pixel clustering and graph analysis
- The methods discussed so far do segmentation but not semantic segmentation (regions with no labels)
- How to get labels?
 - Unlabelled regions can be input to an object classification method
 - Or, segmentation and classification can be done simultaneously (next lecture)